

# Python FastAPI Backend Deployment Guide

## Table of Contents

- 
1. [Pre-Deployment Checklist](#)
  2. [Environment Variables Reference](#)
  3. [Platform-Specific Deployment](#)
    - Railway (Recommended)
    - Render
    - Vercel
  4. [CORS Configuration](#)
  5. [Post-Deployment Steps](#)
  6. [Testing Your Deployment](#)
  7. [Troubleshooting Guide](#)
- 

## Pre-Deployment Checklist

### Files Verification

Before deploying, ensure these files exist in `/home/ubuntu/basketball_app/python-backend/` :

- [x] `app/main.py` - Main FastAPI application
- [x] `app/pose_detection.py` - MediaPipe pose detection logic
- [x] `app/skeleton_drawing.py` - Skeleton overlay rendering
- [x] `app/anatomical_skeleton.py` - AI skeleton generation
- [x] `app/models.py` - Pydantic data models
- [x] `requirements.txt` - Python dependencies
- [x] `Dockerfile` - Container configuration
- [x] `.env.example` - Environment variables template

### Dependencies Verification

Your `requirements.txt` includes:

```
fastapi==0.109.2
uvicorn[standard]==0.27.1
python-multipart==0.0.9
mediapipe==0.10.9
opencv-python==4.9.0.80
Pillow==10.2.0
numpy==1.26.4
replicate==0.25.1
requests==2.31.0
pydantic==2.6.1
python-dotenv==1.0.1
starlette==0.36.3
```

## Key Dependencies:

- **MediaPipe**: For real-time pose detection
- **OpenCV**: Image processing
- **Replicate**: AI-powered anatomical skeleton generation
- **FastAPI + Uvicorn**: Web framework and ASGI server

## 🔑 Required Environment Variables

All platforms require these environment variables:

Variable	Required	Default	Description
REPLICATE_API_TOKEN	✓ Yes	-	API token for AI skeleton generation
ALLOWED_ORIGINS	✓ Yes	-	Comma-separated list of allowed CORS origins
HOST	⚠ Platform-specific	0.0.0.0	Server host (usually 0.0.0.0)
PORT	⚠ Platform-specific	8000	Server port (Railway uses dynamic \$PORT)
ME-DIPIPE_MODEL_COMPLEXITY	Optional	2	Pose detection accuracy (0-2, higher = more accurate)

## 🔒 Environment Variables Reference

### 1. REPLICATE\_API\_TOKEN

**Value:** r8\_XVbSqNpDmahHdfRWdjmiivN2ZNPk3MUH2w1N4x

**Purpose:** Authenticates with Replicate API for AI-powered anatomical skeleton generation.

#### Where it's used:

- /ai-skeleton endpoint for detailed anatomical overlays
- Processes video frames with Replicate's ML models

⚠ **Security Note:** Never commit this token to version control!

### 2. ALLOWED\_ORIGINS

**Format:** Comma-separated URLs (no spaces)

#### Example:

```
ALLOWED_ORIGINS=https://your-abacus-frontend.vercel.app,http://localhost:3000,http://localhost:3001
```

**Purpose:** Controls which domains can make requests to your API (CORS security).

**⚠ Critical:** You MUST include your deployed Abacus AI frontend URL here!

#### Finding Your Frontend URL:

1. Your frontend is deployed on Abacus AI
2. The URL format is typically: `https://[app-name].abacus.ai` or similar
3. Check your Abacus AI deployment dashboard for the exact URL
4. Add this URL to `ALLOWED_ORIGINS`

#### Common Patterns:

```
# For development + production
ALLOWED_ORIGINS=https://basketball-app.abacus.ai,http://localhost:3000

# With wildcards (use cautiously)
ALLOWED_ORIGINS=https://*.abacus.ai,http://localhost:3000

# Multiple environments
ALLOWED_ORIGINS=https://basketball-prod.abacus.ai,https://basketball-staging.abacus.ai,http://localhost:3000
```

## 3. HOST

**Value:** `0.0.0.0`

**Purpose:** Binds the server to all network interfaces, allowing external connections.

#### Platform Notes:

- Railway: Use `0.0.0.0`
- Render: Use `0.0.0.0`
- Vercel: Not applicable (serverless)

## 4. PORT

**Default:** `8000`

#### Platform-Specific Values:

- **Railway:** Use `$PORT` (Railway provides this dynamically)
- **Render:** Use `10000` or `$PORT`
- **Vercel:** Not applicable (serverless)

**Important:** Railway automatically injects `$PORT`, so set it as `PORT=$PORT` or let Railway handle it.

## 5. MEDIAPIPE\_MODEL\_COMPLEXITY

**Value:** `0, 1, or 2`

**Default:** `2` (highest accuracy)

**Purpose:** Controls MediaPipe pose detection model complexity.

Value	Performance	Accuracy	Use Case
0	Fastest	Lower	Real-time, low-latency
1	Balanced	Medium	General use
2	Slower	Highest	Production, detailed analysis

**Recommendation:** Keep at 2 for production unless experiencing performance issues.

---

## Platform-Specific Deployment

### Railway Deployment (⭐ RECOMMENDED)

#### Why Railway?

- Best support for Python + Docker
- Automatic HTTPS
- Simple environment variable management
- Built-in CI/CD from GitHub
- Free tier available (\$5 credit/month)
- Excellent MediaPipe/OpenCV support

#### Step-by-Step Instructions

##### 1. Prepare Your Repository

###### Option A: Use Existing GitHub Repository

Your code is already in: <https://github.com/baller70/BasketballAnalysisAssessmentApp.git>

Ensure the backend is in a clear directory structure:

```
BasketballAnalysisAssessmentApp/
├── python-backend/
│   ├── app/
│   ├── Dockerfile
│   ├── requirements.txt
│   └── .env.example
```

###### Option B: Create a New Repository (Backend Only)

```
cd /home/ubuntu/basketball_app/python-backend
git init
git add .
git commit -m "Initial backend commit"
git remote add origin https://github.com/yourusername/basketball-backend.git
git push -u origin main
```

## 2. Sign Up for Railway

1. Go to [railway.app](https://railway.app) (<https://railway.app>)
2. Click “Start a New Project”
3. Sign in with GitHub (recommended for auto-deployment)

### Screenshot Description:

- Railway homepage with purple gradient background
- “Start a New Project” button in center
- GitHub, GitLab, and email sign-in options

## 3. Create New Project

1. Click “Deploy from GitHub repo”
2. Authorize Railway to access your repositories
3. Select `BasketballAnalysisAssessmentApp` repository
4. Railway will detect the Dockerfile automatically

### Screenshot Description:

- Repository selection screen
- List of your GitHub repositories
- Search bar at top
- “Deploy” button next to each repo

## 4. Configure Root Directory (Important!)

Since your backend is in a subdirectory:

1. Click on the deployed service
2. Go to “Settings” tab
3. Find “Root Directory” setting
4. Set it to: `python-backend`
5. Click “Save”

### Screenshot Description:

- Settings page with tabs (General, Variables, Networking, etc.)
- Root Directory field highlighted
- Text input showing “python-backend”
- Blue “Save” button

## 5. Configure Environment Variables

1. Click “Variables” tab
2. Click “+ New Variable”
3. Add each variable:

### Variable 1: REPLICATE\_API\_TOKEN

Name: REPLICATE_API_TOKEN
Value: r8_XVbSqNpDmahHdfRWdjivN2ZNPk3MUH2w1N4x

### Variable 2: ALLOWED\_ORIGINS

Name: ALLOWED_ORIGINS
Value: https://your-abacus-frontend-url.com,http://localhost:3000

**⚠️ Replace `your-abacus-frontend-url.com` with your actual Abacus AI frontend URL!**

### Variable 3: HOST

Name: HOST
Value: 0.0.0.0

### Variable 4: PORT

Name: PORT
Value: \$PORT

**⚠️ Important:** Railway automatically provides `$PORT`, this ensures your app uses it.

### Variable 5: MEDIPIPE\_MODEL\_COMPLEXITY

Name: MEDIPIPE_MODEL_COMPLEXITY
Value: 2

#### Screenshot Description:

- Variables page with list of environment variables
- Each variable shown as a row with Name and Value columns
- “+ New Variable” button at top
- Eye icon to show/hide sensitive values

## 6. Enable Public Networking

1. Go to “Settings” > “Networking”
2. Click “Generate Domain”
3. Railway will create a public URL like: [https://i.ytimg.com/vi/zX-\\_NftC1K0/maxresdefault.jpg](https://i.ytimg.com/vi/zX-_NftC1K0/maxresdefault.jpg)
4. **Copy this URL** - you’ll need it later!

#### Screenshot Description:

- Networking settings page
- “Generate Domain” button
- Generated domain shown in blue text
- Copy icon next to the domain

## 7. Deploy!

Railway automatically deploys when you:

- Push to GitHub (if connected)
- Change environment variables
- Click “Deploy” manually

### **Monitor deployment:**

1. Go to “**Deployments**” tab
2. Click on the latest deployment
3. View real-time logs
4. Wait for “**Build successful**” and “**Deployment live**”

### **Screenshot Description:**

- Deployments page showing build logs
- Green “Deployment live” status
- Logs showing unicorn server starting
- Timestamp for each log entry

## **8. Verify Deployment**

Test the `/health` endpoint:

```
curl https://your-railway-url.up.railway.app/health
```

### **Expected Response:**

```
{
  "status": "healthy",
  "version": "1.0.0",
  "mediapipe_available": true
}
```

## **Railway Troubleshooting**

### **Problem: Build fails with “No Dockerfile found”**

- Solution: Verify “Root Directory” is set to `python-backend`

### **Problem: Port binding error**

- Solution: Ensure `PORT=$PORT` in environment variables

### **Problem: MediaPipe import errors**

- Solution: Railway’s Docker support handles this automatically via your Dockerfile

## **Render Deployment**

### **Why Render?**

- Free tier available
- Good Docker support
- Automatic HTTPS
- Slower cold starts on free tier
- Free tier sleeps after inactivity

## **Step-by-Step Instructions**

### **1. Sign Up for Render**

1. Go to [render.com](https://render.com) (`https://render.com`)
2. Click “**Get Started**”

3. Sign up with GitHub (recommended)

**Screenshot Description:**

- Render homepage with blue/purple gradient
- “Get Started” button in navigation
- Signup options including GitHub

## 2. Create New Web Service

1. Click “**New +**” button (top right)
2. Select “**Web Service**”
3. Connect your GitHub repository
4. Select `BasketballAnalysisAssessmentApp`

**Screenshot Description:**

- Dashboard with “New +” dropdown menu
- Options: Web Service, Static Site, PostgreSQL, etc.
- Repository connection screen

## 3. Configure Service Settings

**Basic Settings:**

- **Name:** `basketball-backend`
- **Region:** Choose closest to your users
- **Branch:** `main`
- **Root Directory:** `python-backend`
- **Runtime:** `Docker`

**Screenshot Description:**

- Service configuration form
- Dropdown menus for region and runtime
- Text fields for name and root directory

## 4. Configure Build Settings

**Build Command:** (Leave empty - Docker handles this)

**Start Command:**

```
uvicorn app.main:app --host 0.0.0.0 --port $PORT
```

Render automatically injects `$PORT` environment variable.

**Screenshot Description:**

- Build settings section
- “Build Command” text area (empty)
- “Start Command” text area with uvicorn command

## 5. Choose Plan

- **Free:** \$0/month (sleeps after 15min inactivity, 750 hours/month)
- **Starter:** \$7/month (always on, better performance)

**Recommendation:** Start with Free, upgrade if needed.

**Screenshot Description:**

- Pricing tiers displayed as cards

- Free tier highlighted
- Feature comparison table

## 6. Add Environment Variables

Before clicking “Create Web Service”:

1. Scroll to “**Environment Variables**” section
2. Click “**Add Environment Variable**”
3. Add each variable:

### Variables to Add:

```
REPLICATE_API_TOKEN=r8_XVbSqNpDmahHdfRWDjmiivN2ZNPk3MUH2w1N4x
ALLOWED_ORIGINS=https://your-abacus-frontend-url.com,http://localhost:3000
HOST=0.0.0.0
PORT=10000
MEDIAPIPE_MODEL_COMPLEXITY=2
```

**⚠ Replace `your-abacus-frontend-url.com` with your actual frontend URL!**

### Screenshot Description:

- Environment variables section
- Key-value pairs for each variable
- “Add Environment Variable” button
- Eye icons to show/hide values

## 7. Deploy

1. Click “**Create Web Service**”
2. Render will:
  - Clone your repository
  - Build Docker image
  - Deploy container
  - Assign public URL

**Deployment URL:** `https://basketball-backend.onrender.com` (or similar)

### Screenshot Description:

- Deployment in progress screen
- Real-time build logs
- Progress bar showing build stages
- Green “Live” badge when complete

## 8. Update CORS (If Needed)

After deployment, if you get CORS errors:

1. Go to “**Environment**” tab
2. Edit `ALLOWED_ORIGINS`
3. Add your Render URL: `https://basketball-backend.onrender.com`
4. Click “**Save Changes**” (triggers redeployment)

## 9. Verify Deployment

```
curl https://basketball-backend.onrender.com/health
```

### Expected Response:

```
{
  "status": "healthy",
  "version": "1.0.0",
  "mediapipe_available": true
}
```

## Render Troubleshooting

### Problem: “Application failed to respond” error

- **Solution 1:** Check logs for port binding issues
- **Solution 2:** Ensure `PORT=10000` or `PORT=$PORT`
- **Solution 3:** Verify Dockerfile exposes correct port

### Problem: Build fails with OpenCV errors

- **Solution:** Your Dockerfile already includes required system dependencies:

```
dockerfile
RUN apt-get update && apt-get install -y \
    libgl1-mesa-glx \
    libglib2.0-0 \
    libsm6 \
    libxext6 \
    libxrender-dev
```

### Problem: Free tier sleeps too often

- **Solution 1:** Upgrade to paid tier (\$7/month)
- **Solution 2:** Use a keep-alive service (e.g., UptimeRobot) to ping every 14 minutes
- **Solution 3:** Accept 30-60 second cold start delays

## Vercel Deployment

### ⚠️ Important Limitations

Vercel is **primarily designed for Node.js serverless functions**. Python support exists but has significant limitations:

#### Limitations:

- ✗ No native Docker support
- ✗ Serverless functions have 50MB size limit
- ✗ MediaPipe + OpenCV exceed size limits
- ✗ 10-second execution timeout (free tier)
- ✗ Limited Python package support

**Verdict:** ✗ NOT RECOMMENDED for this backend.

## Alternative: Vercel for Frontend + Railway/Render for Backend

#### Recommended Architecture:

1. Deploy **Next.js frontend** on Vercel (optimal)
2. Deploy **Python backend** on Railway or Render
3. Configure CORS to allow Vercel frontend → Railway/Render backend

## If You Must Try Vercel (Not Recommended)

### Prerequisites

You would need to:

1. Convert FastAPI app to Vercel serverless functions
2. Remove MediaPipe/OpenCV (too large)
3. Use external APIs for pose detection
4. Rewrite significant portions of code

### Steps (High-Level):

#### 1. Install Vercel CLI:

```
npm install -g vercel
```

#### 1. Create `vercel.json`:

```
{
  "builds": [
    {
      "src": "app/main.py",
      "use": "@vercel/python"
    }
  ],
  "routes": [
    {
      "src": "/(.*)",
      "dest": "app/main.py"
    }
  ]
}
```

#### 1. Strip Dependencies:

Remove from `requirements.txt`:

- mediapipe
- opencv-python
- Use lighter alternatives or external APIs

#### 2. Deploy:

```
vercel --prod
```

### Why This Won't Work:

- Your backend relies heavily on MediaPipe and OpenCV
- These libraries are essential for pose detection
- No lightweight alternatives exist that maintain functionality

**Conclusion:** ❌ Stick with Railway or Render for this backend.

## CORS Configuration

### Understanding CORS

#### What is CORS?

Cross-Origin Resource Sharing (CORS) is a security mechanism that controls which domains can access your API.

#### Why It Matters:

Your frontend (Abacus AI) needs permission to call your backend API. Without proper CORS configuration, browsers will block requests.

### Current CORS Setup

Your backend (`app/main.py`) includes:

```
app.add_middleware(
    CORSMiddleware,
    allow_origins=[
        "http://localhost:3000",
        "http://localhost:3001",
        "http://127.0.0.1:3000",
        "http://127.0.0.1:3001",
        "https://*.vercel.app",
    ],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

### You MUST Update ALLOWED\_ORIGINS

**Problem:** Hardcoded origins in code won't work for your deployed Abacus AI frontend.

**Solution:** Use environment variable (already implemented).

#### In your deployment platform:

```
ALLOWED_ORIGINS=https://your-abacus-app.abacus.ai,http://localhost:3000
```

### Dynamic CORS Configuration

Update your `app/main.py` (if needed):

```
import os

# Get allowed origins from environment variable
allowed_origins_str = os.getenv("ALLOWED_ORIGINS", "http://localhost:3000")
allowed_origins = [origin.strip() for origin in allowed_origins_str.split(",")]

app.add_middleware(
    CORSMiddleware,
    allow_origins=allowed_origins,
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

## Security Best Practices

### ✓ DO:

#### 1. Use specific domains:

```
bash
ALLOWED_ORIGINS=https://basketball-app.abacus.ai
```

#### 2. List each domain explicitly:

```
bash
ALLOWED_ORIGINS=https://prod.abacus.ai,https://staging.abacus.ai
```

#### 3. Include localhost for development:

```
bash
ALLOWED_ORIGINS=https://prod.abacus.ai,http://localhost:3000
```

### ✗ DON'T:

#### 1. Avoid wildcard for production:

```
bash
# Insecure!
ALLOWED_ORIGINS=*
```

#### 2. Don't expose sensitive origins:

```
bash
# Only include trusted domains
ALLOWED_ORIGINS=https://your-app.com,https://malicious-site.com # ✗
```

## Testing CORS

### Test from Browser Console:

1. Open your Abacus AI frontend in browser
2. Open Developer Tools (F12)
3. Go to Console tab
4. Run:

```
fetch('https://your-backend-url.up.railway.app/health')
  .then(res => res.json())
  .then(data => console.log(data))
  .catch(err => console.error('CORS Error:', err));
```

### Expected Result (Success):

```
{
  status: "healthy",
  version: "1.0.0",
  mediapipe_available: true
}
```

### CORS Error (Failure):

```
Access to fetch at 'https://...' from origin 'https://...' has been blocked by CORS policy
```

**Solution:** Add your frontend domain to `ALLOWED_ORIGINS` and redeploy.

## ✓ Post-Deployment Steps

### 1. Copy Your Backend URL

**Railway:**

```
https://basketball-backend-production-xxxx.up.railway.app
```

**Render:**

```
https://basketball-backend.onrender.com
```

Copy this URL - you'll need it for the frontend!

### 2. Update Frontend Environment Variable

**Location:** Your Abacus AI frontend deployment settings

**Variable to Update:**

```
NEXT_PUBLIC_PYTHON_API_URL=https://your-backend-url.up.railway.app
```

**Steps:**

1. Go to your Abacus AI deployment dashboard
2. Find environment variables settings
3. Update `NEXT_PUBLIC_PYTHON_API_URL` with your backend URL
4. Redeploy the frontend

⚠ **Important:** Remove trailing slash from URL!

✓ Correct:

```
NEXT_PUBLIC_PYTHON_API_URL=https://backend.railway.app
```

✗ Incorrect:

```
NEXT_PUBLIC_PYTHON_API_URL=https://backend.railway.app/
```

### 3. Update CORS in Backend

Ensure your frontend URL is in `ALLOWED_ORIGINS`:

```
ALLOWED_ORIGINS=https://basketball-app.abacus.ai,http://localhost:3000
```

**On Railway/Render:**

1. Go to Environment Variables
2. Edit `ALLOWED_ORIGINS`
3. Add your Abacus AI frontend URL
4. Save (triggers automatic redeployment)

## 4. Test All Endpoints

See [Testing Your Deployment](#) section below.

### Testing Your Deployment

#### Test 1: Health Check

**Purpose:** Verify server is running and MediaPipe is available.

**Command:**

```
curl https://your-backend-url/health
```

**Expected Response:**

```
{
  "status": "healthy",
  "version": "1.0.0",
  "mediapipe_available": true
}
```

**Troubleshooting:**

-  Connection refused → Check deployment status
-  mediapipe\_available: false → Check Dockerfile dependencies

#### Test 2: Analyze Endpoint (Image Upload)

**Purpose:** Test pose detection with a sample image.

**Command:**

```
curl -X POST https://your-backend-url/analyze \
  -F "file=@/path/to/basketball-image.jpg" \
  -H "Content-Type: multipart/form-data"
```

**Expected Response:**

```
{
  "keypoints": [
    {"x": 0.5, "y": 0.3, "z": -0.1, "visibility": 0.98, "name": "nose"},
    {"x": 0.48, "y": 0.28, "z": -0.05, "visibility": 0.95, "name": "left_eye"},
    ...
  ],
  "confidence": 0.92,
  "shooting_pose_detected": true,
  "image_dimensions": {"width": 1920, "height": 1080}
}
```

**Troubleshooting:**

-  400 Bad Request → Check file format (must be image)
-  500 Internal Server Error → Check logs for MediaPipe errors

## Test 3: AI Skeleton Endpoint

**Purpose:** Test Replicate API integration.

**Command:**

```
curl -X POST https://your-backend-url/ai-skeleton \
-F "file=@/path/to/basketball-image.jpg" \
-F "strength=0.8" \
-F "detail_level=high" \
-H "Content-Type: multipart/form-data"
```

**Expected Response:**

```
{
  "output_image": "https://lh7-rt.googleusercontent.com/docsz/
AD_4nXfj5huwABX38xsU1SNq1yr-PUswy-nVH5Cc6h0InAtdcvAHzXvmRDKAkPQy6JI1U9BD1M3-
foBI56GI_d1ASjzBlJQ03SZJlMctBnCkzdx4QPCjN8Ry_x25znAL19pFriEsww1mVBQ?key=-ICjWANGPGN-
BY6zqswhs_w",
  "processing_time": 3.45
}
```

**Troubleshooting:**

- ✗ 401 Unauthorized → Check `REPLICATE_API_TOKEN`
- ✗ Timeout → Replicate API may be slow, increase timeout

## Test 4: CORS Test (From Browser)

**Purpose:** Verify frontend can access backend.

**Steps:**

1. Open Abacus AI frontend in browser
2. Open Developer Tools (F12) → Console
3. Run:

```
fetch('https://your-backend-url/health')
  .then(res => res.json())
  .then(data => console.log('✓ CORS working:', data))
  .catch(err => console.error('✗ CORS error:', err));
```

**Expected Result:**

```
✓ CORS working: {status: "healthy", version: "1.0.0", ...}
```

**If CORS Error:**

1. Check `ALLOWED_ORIGINS` includes your frontend URL
2. Verify no typos in domain names
3. Ensure no trailing slashes
4. Redeploy backend after changes

## Test 5: Full Integration Test

**Purpose:** Test complete workflow from frontend.

**Steps:**

1. Go to your Abacus AI frontend
2. Navigate to analysis page
3. Upload a basketball shooting image
4. Verify:
  - Image uploads successfully
  - Pose detection completes
  - Skeleton overlay displays
  - Elite shooter comparisons appear

**Expected Behavior:**

- Image preview shows immediately
  - Loading spinner during analysis
  - Skeleton overlay renders with keypoints
  - No CORS errors in console
- 



## Troubleshooting Guide

### Common Issues & Solutions

#### CORS Errors

**Error:**

```
Access to fetch at 'https://backend.railway.app/analyze' from origin 'https://front-end.abacus.ai'
has been blocked by CORS policy: No 'Access-Control-Allow-Origin' header is present
```

**Cause:** Backend doesn't allow requests from your frontend domain.**Solutions:****1. Add frontend URL to ALLOWED\_ORIGINS:**

```
bash
```

```
ALLOWED_ORIGINS=https://frontend.abacus.ai,http://localhost:3000
```

**2. Verify no typos:**

- https://frontend.abacus.ai (correct)
- https://frontend.abacus.ai/ (trailing slash)
- http://frontend.abacus.ai (http vs https)

**3. Check environment variable is loaded:**

- View logs: Should see `allow_origins=[...]` on startup
- Verify Railway/Render environment variables saved correctly

**4. Redeploy after changes:**

- Environment variable changes require redeployment

**Verification:**

```
# Check CORS headers
curl -I https://your-backend-url/health \
-H "Origin: https://your-frontend-url"

# Should include:
# Access-Control-Allow-Origin: https://your-frontend-url
```

## ✖ MediaPipe Import Error

### Error:

```
ModuleNotFoundError: No module named 'mediapipe'
```

**Cause:** Dependencies not installed correctly.

### Solutions:

#### 1. Verify requirements.txt includes mediapipe:

```
txt
mediapipe==0.10.9
```

#### 2. Check Dockerfile installs system dependencies:

```
dockerfile
RUN apt-get update && apt-get install -y \
    libgl1-mesa-glx \
    libglib2.0-0 \
    libsm6 \
    libxext6 \
    libxrender-dev
```

#### 3. Rebuild from scratch:

- Railway: Trigger manual redeploy
- Render: Clear build cache and redeploy

#### 4. Check build logs:

- Look for Successfully installed mediapipe-0.10.9
- Watch for “Failed to build” errors

## ✖ OpenCV Import Error

### Error:

```
ImportError: libGL.so.1: cannot open shared object file
```

**Cause:** Missing OpenCV system dependencies.

### Solution:

Your Dockerfile already includes the fix. If error persists:

### 1. Verify Dockerfile includes:

```
dockerfile
RUN apt-get update && apt-get install -y \
    libgl1-mesa-glx \
    libglib2.0-0
```

### 2. Rebuild Docker image:

```
bash
docker build --no-cache -t basketball-backend .
```

### 3. Alternative: Use opencv-python-headless:

In `requirements.txt` :

```
txt
opencv-python-headless==4.9.0.80
(This removes GUI dependencies)
```

## Port Binding Error

### Error:

```
OSError: [Errno 98] Address already in use
```

**Cause:** Wrong PORT configuration.

### Solutions:

#### Railway:

```
PORTR=$PORT
```

Railway injects dynamic port.

#### Render:

```
PORTR=10000
# or
PORTR=$PORT
```

#### Local Testing:

```
PORTR=8000
```

#### Verification:

Check logs for:

```
INFO:     Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
```

## ✖ Replicate API Errors

### Error 1: 401 Unauthorized

```
{"detail": "Invalid Replicate API token"}
```

#### Solution:

- Verify `REPLICATE_API_TOKEN` is set correctly
- Check for extra spaces or newlines
- Ensure token starts with `r8_`

### Error 2: 429 Rate Limited

```
{"detail": "Replicate API rate limit exceeded"}
```

#### Solution:

- Wait and retry
- Upgrade Replicate plan for higher limits
- Implement caching for repeated requests

### Error 3: Timeout

```
replicate.exceptions.ReplicateError: Request timeout
```

#### Solution:

- Replicate's AI models can be slow (30+ seconds)
- Increase timeout in code:

```
python
import replicate
replicate.Client(api_token=REPLICATE_API_TOKEN, timeout=60)
```

## ✖ Image Upload Fails

### Error:

```
{"detail": "File must be an image"}
```

**Cause:** Invalid file type or corrupted upload.

### Solutions:

#### 1. Check file type:

- Supported: JPG, JPEG, PNG, WebP
- Not supported: GIF, SVG, HEIC

#### 2. Verify Content-Type header:

```
bash
curl -X POST ... -F "file=@image.jpg" \
-H "Content-Type: multipart/form-data"
```

### 3. Check file size:

- Max recommended: 10 MB
- Larger files may timeout

### 4. Test with sample image:

```
```bash
# Download test image
curl -o test.jpg https://dummyimage.com/svga

# Test upload
curl -X POST https://your-backend-url/analyze \
-F "file=@test.jpg"
```

```

---

## Frontend Can't Connect

### Symptoms:

- Network errors in browser console
- “Failed to fetch” errors
- Requests never reach backend

### Solutions:

#### 1. Verify NEXT\_PUBLIC\_PYTHON\_API\_URL:

- Check Abacus AI frontend environment variables
- Should be: `https://your-backend-url`
- No trailing slash!

#### 2. Check backend is running:

```
bash
curl https://your-backend-url/health
```

#### 3. Verify HTTPS:

- Backend URL must use HTTPS (not HTTP)
- Railway and Render provide automatic HTTPS

#### 4. Check CORS (again):

- See CORS troubleshooting section above

#### 5. Test from command line:

```
bash
# If this works but frontend doesn't, it's a CORS issue
curl https://your-backend-url/health
```

---

## Deployment Succeeds but API Returns 404

### Error:

```
{"detail": "Not Found"}
```

**Cause:** Incorrect root directory or routing.

**Solutions:****1. Verify endpoint paths:**

- `https://backend-url/health`
- `https://backend-url/api/health`
- `https://backend-url/python-backend/health`

**2. Check Root Directory setting:**

- Railway: Should be `python-backend`
- Render: Should be `python-backend`

**3. Verify main.py location:**

```
python-backend/
  └── app/
    └── main.py ← Should be here
```

**4. Check start command:**

```
bash
uvicorn app.main:app --host 0.0.0.0 --port $PORT
```

Not:

```
bash
uvicorn main:app # ✗ Wrong
```

**✗ Slow Response Times****Symptoms:**

- Requests take 10+ seconds
- Timeouts on image analysis

**Solutions:****1. Reduce MediaPipe complexity:**

```
bash
MEDIPIPE_MODEL_COMPLEXITY=1 # Instead of 2
```

**2. Optimize image size:**

- Resize images before upload
- Recommended: 1920x1080 or smaller

**3. Upgrade hosting plan:**

- Railway: Hobby plan (\$5/month)
- Render: Starter plan (\$7/month)
- Free tiers have limited CPU/memory

**4. Add caching:**

- Cache pose detection results
- Avoid re-processing same images

**5. Use async endpoints:**

- Already implemented with `async def`
- Ensure Uvicorn uses multiple workers:

```
bash
uvicorn app.main:app --workers 4
```

## ✖ Memory Errors

### Error:

```
MemoryError: Unable to allocate array
```

**Cause:** Large images or limited RAM.

### Solutions:

#### 1. Resize images before processing:

Add to `pose_detection.py`:

```
python
max_dimension = 1920
if image.shape[0] > max_dimension or image.shape[1] > max_dimension:
    scale = max_dimension / max(image.shape[:2])
    image = cv2.resize(image, None, fx=scale, fy=scale)
```

#### 2. Upgrade hosting plan:

- Free tiers: 512 MB RAM
- Paid tiers: 2+ GB RAM

#### 3. Reduce model complexity:

```
bash
MEDIPIPE_MODEL_COMPLEXITY=0
```

## Getting Help

### If issues persist:

#### 1. Check deployment logs:

- Railway: Deployments tab → View Logs
- Render: Logs tab (left sidebar)

#### 2. Enable debug mode:

Add to environment variables:

```
bash
LOG_LEVEL=DEBUG
```

#### 3. Test locally:

```
bash
cd /home/ubuntu/basketball_app/python-backend
pip install -r requirements.txt
uvicorn app.main:app --reload
```

#### 4. Contact support:

- Railway: [railway.app/help](https://railway.app/help) (<https://railway.app/help>)
- Render: [render.com/support](https://render.com/support) (<https://render.com/support>)



## Deployment Checklist

---

### Pre-Deployment

- [ ] All files present in `/home/ubuntu/basketball_app/python-backend/`
- [ ] `requirements.txt` includes all dependencies
- [ ] `Dockerfile` configured correctly
- [ ] `.env.example` documented
- [ ] Code pushed to GitHub

### During Deployment

- [ ] Platform account created (Railway/Render)
- [ ] Repository connected
- [ ] Root directory set to `python-backend`
- [ ] Environment variables configured:
- [ ] `REPLICATE_API_TOKEN`
- [ ] `ALLOWED_ORIGINS` (with frontend URL)
- [ ] `HOST=0.0.0.0`
- [ ] `PORT` (platform-specific)
- [ ] `MEDIPIPE_MODEL_COMPLEXITY=2`
- [ ] Public domain generated

### Post-Deployment

- [ ] `/health` endpoint returns 200 OK
- [ ] `mediapipe_available: true` in health check
- [ ] Frontend environment variable updated:
- [ ] `NEXT_PUBLIC_PYTHON_API_URL` set to backend URL
- [ ] Frontend redeployed
- [ ] CORS test passes from browser
- [ ] `/analyze` endpoint tested with sample image
- [ ] `/ai-skeleton` endpoint tested
- [ ] Full integration test completed

### Monitoring

- [ ] Set up uptime monitoring (optional)
- [ ] Configure error alerts
- [ ] Monitor API usage/costs



## Success!

---

If all tests pass, your Python FastAPI backend is successfully deployed and connected to your Abacus AI frontend!

### Next Steps:

1. Share the frontend URL with users
2. Monitor performance and errors
3. Scale as needed

4. Consider adding:

- Rate limiting
- API authentication
- Response caching
- Monitoring/analytics

**Useful Links:**

- Frontend (Abacus AI): <https://your-frontend-url>
  - Backend API: <https://your-backend-url>
  - API Documentation: <https://your-backend-url/docs> (FastAPI auto-generated)
  - Backend Health: <https://your-backend-url/health>
- 



## Additional Resources

- [FastAPI Documentation](https://fastapi.tiangolo.com/) (<https://fastapi.tiangolo.com/>)
  - [Railway Documentation](https://docs.railway.app/) (<https://docs.railway.app/>)
  - [Render Documentation](https://render.com/docs) (<https://render.com/docs>)
  - [MediaPipe Pose Guide](https://google.github.io/mediapipe/solutions/pose.html) (<https://google.github.io/mediapipe/solutions/pose.html>)
  - [Replicate API Docs](https://replicate.com/docs) (<https://replicate.com/docs>)
- 

**Document Version:** 1.0.0

**Last Updated:** December 4, 2025

**Maintainer:** Basketball Analysis Team